

TEMPORAL DATABASES: THEORY AND POSTGRES 2024

Paul A. Jungwirth
15 October 2024
PostgreSQL Chicago

THE PROBLEM

LOST INFORMATION

THE PROBLEM

LOST INFORMATION

- finance: accounting, market data, etc.

THE PROBLEM

LOST INFORMATION

- finance: accounting, market data, etc.
- questionnaires: changing questions, options

THE PROBLEM

LOST INFORMATION

- finance: accounting, market data, etc.
- questionnaires: changing questions, options
- e-commerce: product price, other attributes

THE PROBLEM

LOST INFORMATION

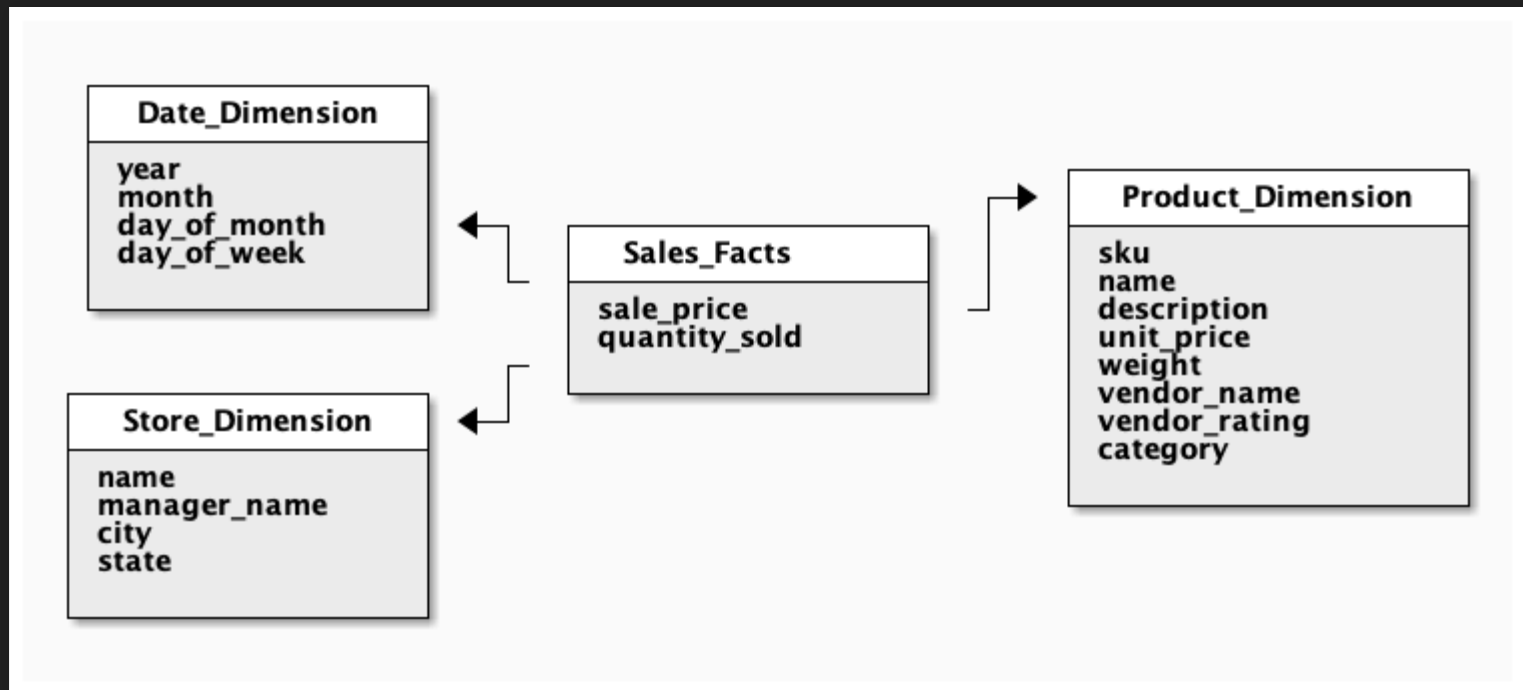
- finance: accounting, market data, etc.
- questionnaires: changing questions, options
- e-commerce: product price, other attributes
- real estate: house renovations

THE PROBLEM

LOST INFORMATION

- finance: accounting, market data, etc.
- questionnaires: changing questions, options
- e-commerce: product price, other attributes
- real estate: house renovations
- employees: position, salary, employment period

OLAP PROBLEMS TOO



"SLOWLY-CHANGING DIMENSIONS"

- Type I: Overwrite it
- Type II: Add a Row
- Type III: Add a Column (good for one change)

"SLOWLY-CHANGING DIMENSIONS"

- Type I: Overwrite it
- Type II: Add a Row
- Type III: Add a Column (good for one change)

later:

- Type IV: Mini-dimensions
- Type V: Mini-dimensions with outriggers
- Type VI: Original vs Current
- Type VII: Type I + Type II

THE FIRST DBA?



of access paths in the total model for the community of users of a data bank would eventually become excessively large.


1.3. A RELATIONAL VIEW OF DATA

The term *relation* is used here in its accepted mathematical sense. Given sets S_1, S_2, \dots, S_n (not necessarily distinct), R is a relation on these n sets if it is a set of n -tuples each of which has its first element from S_1 , its second element from S_2 , and so on.¹ We shall refer to S_j as the j th domain of R . As defined above, R is said to have degree n . Relations of degree 1 are often called *unary*, degree 2 *binary*, degree 3 *ternary*, and degree n *n-ary*.


For expository reasons, we shall frequently make use of an array representation of relations, but it must be remembered that this particular representation is not an essential part of the relational view being expounded. An ar-

¹ More concisely, R is a subset of the Cartesian product $S_1 \times S_2 \times \dots \times S_n$.


RESEARCH



RESEARCH



RESEARCH



TEMPORAL IS DISTINCT FROM TIME-SERIES

time-series	temporal
single timestamp	two timestamps
records events	records things, "versions"
IoT sensors, finance	auditing, history
challenge is scale	challenge is complexity
partitioning	ranges, exclusion constraints
TimescaleDB	periods, pg_bitemporal

TWO DIMENSIONS

Application Time	System Time
history of the thing	history of the database
application features	auditing, compliance
user can edit	immutable
maintained by you	maintained by Postgres
constraints matter	look Ma, no hands!
periods, pg_bitemporal	temporal_tables, pg_audit_log
nothing	Rails: papertrail, audited, chronomodel, ...

TERMINOLOGY

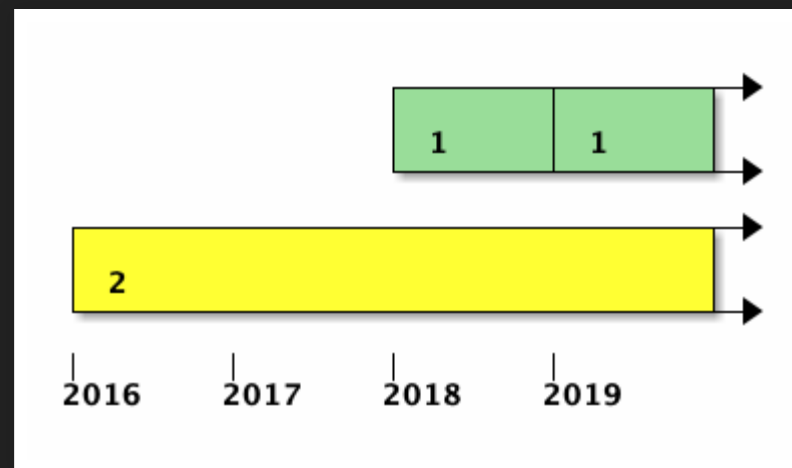
Snodgrass	valid time	transaction time
Fowler	actual time	record time
Date/Darwen/Lorentzos	stated time	logged time
Johnston	effective time/ state time	assertion time
SQL:2011	application time	system time

TEMPORAL EXAMPLE

products				
id	name	price	valid_from	valid_til
1	shoe	\$5	Jan 2018	Jan 2019
1	shoe	\$7	Jan 2019	
2	snow	\$2	Jan 2016	


TEMPORAL EXAMPLE

products				
id	name	price	valid_from	valid_til
1	shoe	\$5	Jan 2018	Jan 2019
1	shoe	\$7	Jan 2019	
2	snow	\$2	Jan 2016	



TEMPORAL EXAMPLE

products			
id	name	price	valid_at
1	shoe	\$5	[Jan 2018, Jan 2019) [Jan 2019,)
1	shoe	\$7	
2	snow	\$2	[Jan 2016,)



RANGE OPERATORS

Operator	Description	Example	Result
=	equal	<code>int4range(1,5) = '[1,4]':int4range</code>	t
<>	not equal	<code>numrange(1.1,2.2) <> numrange(1.1,2.3)</code>	t
<	less than	<code>int4range(1,10) < int4range(2,3)</code>	t
>	greater than	<code>int4range(1,10) > int4range(1,5)</code>	t
<=	less than or equal	<code>numrange(1.1,2.2) <= numrange(1.1,2.2)</code>	t
>=	greater than or equal	<code>numrange(1.1,2.2) >= numrange(1.1,2.0)</code>	t

MORE OPERATORS

@>	contains range	int4range(2,4) @> int4range(2,3)	t
@>	contains element	'[2011-01-01,2011-03-01)>::tsrange @> '2011-01-10'::timestamp	t
<@	range is contained by	int4range(2,4) <@ int4range(1,7)	t
<@	element is contained by	42 <@ int4range(1,7)	f
&&	overlap (have points in common)	int8range(3,7) && int8range(4,12)	t
<<	strictly left of	int8range(1,10) << int8range(100,110)	t
>>	strictly right of	int8range(50,60) >> int8range(20,30)	t

AND MORE

&<	does not extend to the right of	<code>int8range(1,20) &< int8range(18,20)</code>	<code>t</code>
&>	does not extend to the left of	<code>int8range(7,20) &> int8range(5,10)</code>	<code>t</code>
- -	is adjacent to	<code>numrange(1.1,2.2) - - numrange(2.2,3.3)</code>	<code>t</code>
+	union	<code>numrange(5,15) + numrange(10,20)</code>	<code>[5,20)</code>
*	intersection	<code>int8range(5,15) * int8range(10,20)</code>	<code>[10,15)</code>
-	difference	<code>int8range(5,15) - int8range(10,20)</code>	<code>[5,10)</code>

NON-UNIQUE PKS

products			
id	name	price	valid_at
1 1	shoe shoe	\$5 \$7	[Jan 2018, Jan 2019) [Jan 2019,)
2	snow	\$2	[Jan 2016,)
3 3	sail sail	\$8 \$9	[Jan 2016,) [Jan 2017, Jan 2018)

EXCLUSION CONSTRAINTS

```
ALTER TABLE products  
ADD CONSTRAINT products_pk  
EXCLUDE USING gist  
(id WITH =, valid_at WITH &&);
```

TEMPORAL PKS

```
ALTER TABLE products  
ADD CONSTRAINT products_pk  
PRIMARY KEY (id, valid_at WITHOUT OVERLAPS);
```

TEMPORAL PKS

```
ALTER TABLE products  
ADD CONSTRAINT products_pk  
PRIMARY KEY (id, valid_at WITHOUT OVERLAPS);
```

```
ALTER TABLE products  
ADD CONSTRAINT products_uq  
UNIQUE (id, valid_at WITHOUT OVERLAPS);
```

FOREIGN KEYS

products		
id	name	price
1	shoe	\$5
2	snow	\$2



variants		
id	product_id	size
1 2	1 1	5 8
3	2	5
4	3	1

TEMPORAL FOREIGN KEYS

products			
id	name	price	valid_at
1	shoe	\$5	[Jan 2018, Jan 2019)
1	shoe	\$7	[Jan 2019,)
2	snow	\$2	[Jan 2016,)



variants			
id	product_id	size	valid_at
1	1	5	[Jan 2018, Mar 2018)
2	1	8	[Jan 2018, Jan 2020)
3	2	5	[Jan 2014, Jan 2015)

TEMPORAL FOREIGN KEYS

```
CHECK (  
  NOT EXISTS (  
    SELECT 1  
    FROM    variants AS v  
    -- There was a p when v started:  
    WHERE NOT EXISTS (  
      SELECT 1  
      FROM    products AS p  
      WHERE   v.product_id = p.id  
      AND     coalesce(lower(p.valid_at), '-infinity')  
              <= coalesce(lower(v.valid_at), '-infinity')  
      AND     coalesce(lower(v.valid_at), '-infinity')  
              <  coalesce(upper(p.valid_at), 'infinity'))  
    -- ...  
  )  
)
```

TEMPORAL FOREIGN KEYS

```
-- ...
-- There was a p when v ended:
OR NOT EXISTS (
  SELECT 1
  FROM    products AS p
  WHERE   v.product_id = p.id
  AND     coalesce(lower(p.valid_at), '-infinity')
          < coalesce(upper(v.valid_at), 'infinity')
  AND     coalesce(upper(v.valid_at), 'infinity')
          <= coalesce(upper(p.valid_at), 'infinity'))
-- ...
```

TEMPORAL FOREIGN KEYS

```
-- ...
-- There are no gaps in p throughout v:
OR EXISTS (
  SELECT 1
  FROM    products AS p
  WHERE   v.product_id = p.id
  AND     coalesce(lower(v.valid_at), '-infinity')
          < coalesce(upper(p.valid_at), 'infinity')
  AND     coalesce(upper(p.valid_at), 'infinity')
          < coalesce(upper(v.valid_at), 'infinity')
-- ...
```


TEMPORAL FOREIGN KEYS

```
-- ...  
AND NOT EXISTS (  
    SELECT 1  
    FROM    products AS p2  
    WHERE   p2.id = p.id  
    AND     coalesce(lower(p2.valid_at), '-infinity')  
            <= coalesce(upper(p.valid_at), 'infinity')  
    AND     coalesce(upper(p.valid_at), 'infinity')  
            < coalesce(upper(p2.valid_at), 'infinity'))))
```

TEMPORAL FOREIGN KEYS




```
ALTER TABLE variants
ADD CONSTRAINT variants_product_id_fk
FOREIGN KEY (product_id, PERIOD valid_at)
REFERENCES (id, PERIOD valid_at);
```


QUERIES

snapshot ("current")	at a given moment	returns a traditional table (removes <code>valid_at</code>)	<code>WHERE valid_at @> t</code>
sequenced	across time	returns a temporal table (preserves <code>valid_at</code>)	nothing, or <code>WHERE valid_at && r</code>
non- sequenced	time is just another column	returns ???	


TEMPORAL INSERT




TEMPORAL INSERT



TEMPORAL INSERT




TEMPORAL INSERT




TEMPORAL UPDATE

TEMPORAL UPDATE





TEMPORAL UPDATE




```
UPDATE products
FOR PORTION OF valid_at
    FROM '2018-07-01'
    TO '2019-07-01'
SET price = 9
WHERE id = 'shoe';
```

TEMPORAL UPDATE




TEMPORAL DELETE

TEMPORAL DELETE





TEMPORAL DELETE




```
DELETE FROM products
FOR PORTION OF valid_at
      FROM '2018-07-01'
      TO '2019-07-01'
WHERE   id = 'shoe';
```



TEMPORAL DELETE





TEMPORAL UPSERT




TEMPORAL UPSERT



TEMPORAL UPSERT



SQL:2011



RANGE

```
CREATE TABLE products (  
  id          integer,  
  valid_at   tstzrange,  
  
  name       text,  
  price      decimal(10,2),  
  
  CONSTRAINT pk_products  
    PRIMARY KEY  
    (id, valid_at WITHOUT OVERLAPS)  
);
```

RANGE

```
CREATE TABLE products (  
  id          integer,  
  valid_at    tstzrange,  
  
  name        text,  
  price       decimal(10,2),  
  
  CONSTRAINT pk_products  
    PRIMARY KEY  
    (id, valid_at WITHOUT OVERLAPS)  
);
```

PERIOD

```
CREATE TABLE products (  
  id          integer,  
  valid_from  timestampz NOT NULL,  
  valid_til   timestampz NOT NULL,  
  
  name        text,  
  price       decimal(10,2),  
  
  PERIOD FOR valid_at  
    (valid_from, valid_til),  
  CONSTRAINT pk_products  
    PRIMARY KEY  
    (id, valid_at WITHOUT OVERLAPS)  
);
```

SYSTEM TIME

```
CREATE TABLE products (  
  id          integer,  
  sys_from    timestamp GENERATED ALWAYS AS ROW START,  
  sys_til     timestamp GENERATED ALWAYS AS ROW END,  
  
  name        text,  
  price       decimal(10,2),  
  
  PERIOD FOR SYSTEM_TIME  
    (sys_from, sys_til)  
) WITH SYSTEM VERSIONING;
```

SYSTEM TIME

```
SELECT *  
FROM   products  
FOR SYSTEM_TIME AS OF t;
```

```
SELECT *  
FROM   products  
FOR SYSTEM_TIME FROM t1 TO t2;
```


INNER JOINS

```
SELECT  e.name, e.salary, p.name,  
        e.valid_at * p.valid_at  
FROM    employees AS e  
JOIN    positions AS p  
ON      p.employee_id = e.id  
AND     p.valid_at && e.valid_at
```

OUTER JOINS

offers		
house_id	price	valid_at
1	\$100	[Feb 11, Feb 14)
1	\$150	[Feb 14, Feb 17)
1	\$100	[Feb 17, Feb 22)



reservations		
house_id	customer_id	valid_at
1	1	[Feb 13, Feb 15)
1	2	[Feb 16, Feb 19)

OUTER JOINS

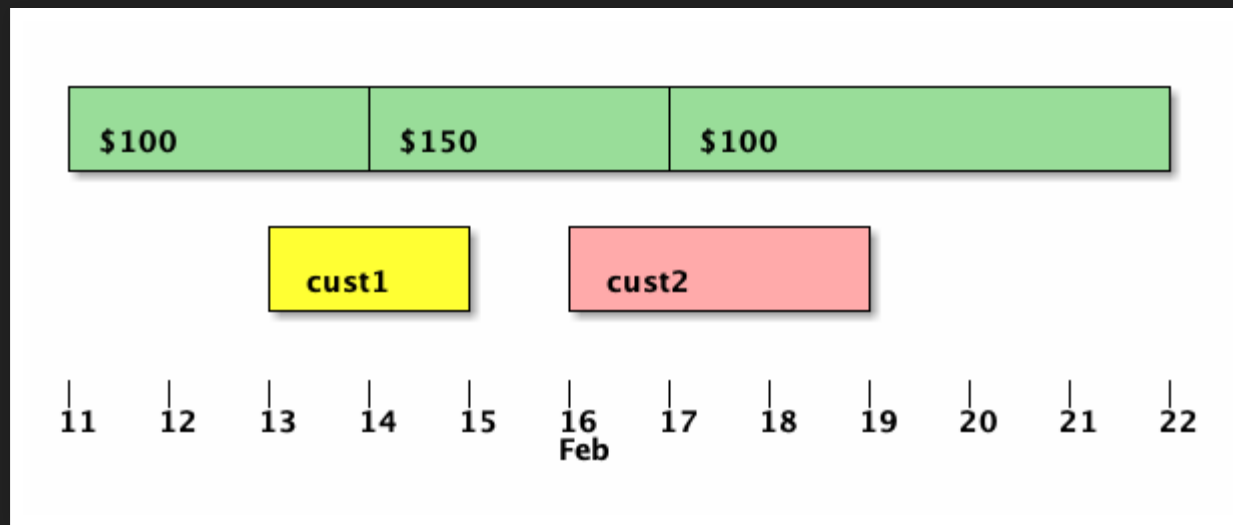
offers		
house_id	price	valid_at
1	\$100	[Feb 11, Feb 14)
1	\$150	[Feb 14, Feb 17)
1	\$100	[Feb 17, Feb 22)



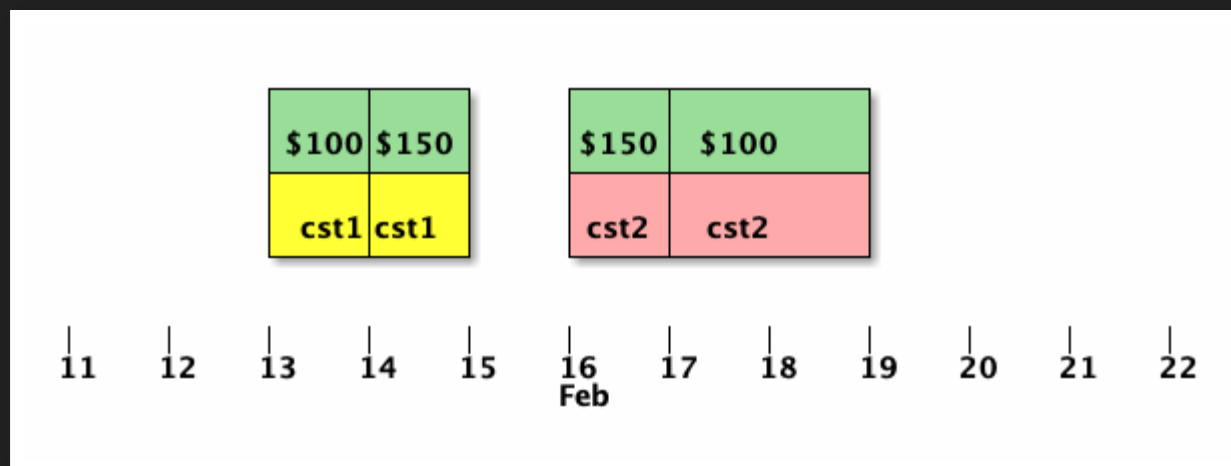
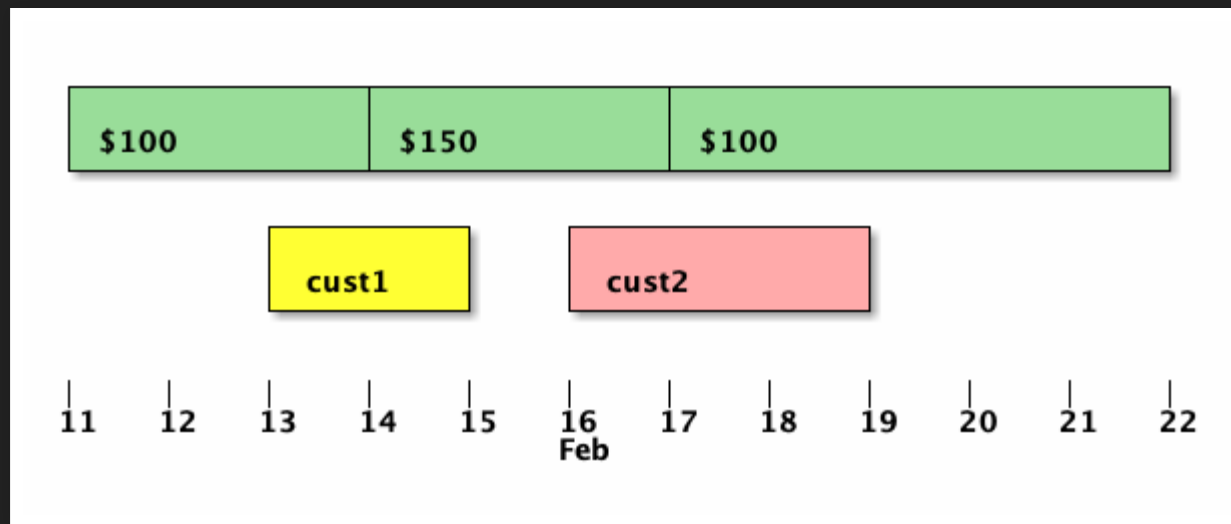
reservations		
house_id	customer_id	valid_at
1	1	[Feb 13, Feb 15)
1	2	[Feb 16, Feb 19)

customer_id	price	valid_at
1	\$100	[Feb 13, Feb 14)
1	\$150	[Feb 14, Feb 15)
2	\$150	[Feb 16, Feb 17)
2	\$100	[Feb 17, Feb 19)

OUTER JOINS



OUTER JOINS



SEMIJOINS

```
SELECT  a.id, a.valid_at
FROM    a
WHERE   EXISTS (
    SELECT  1
    FROM    b
    WHERE   a.id = b.id
    AND     a.valid_at && b.valid_at);
```

SEMIJOINS

```
SELECT  a.id, a.valid_at
FROM    a
WHERE   EXISTS (
    SELECT  1
    FROM    b
    WHERE   a.id = b.id
    AND     a.valid_at && b.valid_at);
```

SEMIJOINS

```
SELECT  a.id,  
        UNNEST(multirange(a.valid_at) * j.valid_at) AS valid_a  
FROM    a  
JOIN (    
    SELECT  b.id, range_agg(b.valid_at) AS valid_at  
    FROM    b  
    GROUP BY b.id  
  ) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at;
```


SEMIJOINS

```
SELECT  a.id,  
        UNNEST(multirange(a.valid_at) * j.valid_at) AS valid_a  
FROM    a  
JOIN (    
    SELECT  b.id, range_agg(b.valid_at) AS valid_at  
    FROM    b  
    GROUP BY b.id  
  ) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at;
```

SEMIJOINS

```
SELECT  a.id,  
        UNNEST(multirange(a.valid_at) * j.valid_at) AS valid_a  
FROM    a  
JOIN (    
    SELECT  b.id, range_agg(b.valid_at) AS valid_at  
    FROM    b  
    GROUP BY b.id  
  ) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at;
```

ANTIJOINS

```
SELECT  a.id,  
        UNNEST(  
            CASE WHEN j.valid_at IS NULL  
                 THEN multirange(a.valid_at)  
                 ELSE multirange(a.valid_at) - j.valid_at END  
        ) AS valid_at  
FROM    a  
LEFT JOIN (  
    SELECT b.id, range_agg(b.valid_at) AS valid_at  
    FROM   b  
    GROUP BY b.id  
) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at  
WHERE NOT isempty(a.valid_at);
```

ANTIJOINS

```
SELECT  a.id,  
        UNNEST(  
            CASE WHEN j.valid_at IS NULL  
                  THEN multirange(a.valid_at)  
                  ELSE multirange(a.valid_at) - j.valid_at END  
        ) AS valid_at  
FROM    a  
LEFT JOIN (  
    SELECT b.id, range_agg(b.valid_at) AS valid_at  
    FROM   b  
    GROUP BY b.id  
) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at  
WHERE NOT isempty(a.valid_at);
```

ANTIJOINS

```
SELECT  a.id,  
        UNNEST(  
            CASE WHEN j.valid_at IS NULL  
                  THEN multirange(a.valid_at)  
                  ELSE multirange(a.valid_at) - j.valid_at END  
        ) AS valid_at  
FROM    a  
LEFT JOIN (  
    SELECT b.id, range_agg(b.valid_at) AS valid_at  
    FROM   b  
    GROUP BY b.id  
) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at  
WHERE NOT isempty(a.valid_at);
```

ANTIJOINS

```
SELECT  a.id,  
        UNNEST(  
            CASE WHEN j.valid_at IS NULL  
                  THEN multirange(a.valid_at)  
                  ELSE multirange(a.valid_at) - j.valid_at END  
        ) AS valid_at  
FROM    a  
LEFT JOIN (  
    SELECT b.id, range_agg(b.valid_at) AS valid_at  
    FROM   b  
    GROUP BY b.id  
) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at  
WHERE NOT isempty(a.valid_at);
```

ANTIJOINS

```
SELECT  a.id,  
        UNNEST(  
            CASE WHEN j.valid_at IS NULL  
                  THEN multirange(a.valid_at)  
                  ELSE multirange(a.valid_at) - j.valid_at END  
        ) AS valid_at  
FROM    a  
LEFT JOIN (  
    SELECT b.id, range_agg(b.valid_at) AS valid_at  
    FROM   b  
    GROUP BY b.id  
) AS j  
ON a.id = j.id AND a.valid_at && j.valid_at  
WHERE NOT isempty(a.valid_at);
```

AGGREGATES

TODO

UNION, INTERSECT, EXCEPT

TODO

MORE

- History UX
- CRUD: REST, GraphQL
- ORM

THANKS!

ME

- <https://github.com/pjungwir/temporal-databases-postgres-talk>
- <https://illuminatedcomputing.com/posts/2017/12/temporal-databases-bibliography/>

RESEARCH

- <https://www2.cs.arizona.edu/~rts/publications.html>
- <http://citeseerx.ist.psu.edu/viewdoc/download;jsessionid=F78723B857463955C76E540DCAB8FDF5?doi=10.1.1.116.7598&rep=rep1&type=pdf>
- <https://files.ifi.uzh.ch/boehlen/Papers/modf174-dignoes.pdf>
- http://www.zora.uzh.ch/id/eprint/130374/1/Extending_the_kernel.pdf
- <https://www.red-gate.com/simple-talk/databases/postgresql/making-temporal-databases-work-part-2-computing-aggregates-across-temporal-versions/>
- https://github.com/pjungwir/temporal_ops

SQL:2011

- <https://www.wiscorp.com/SQLStandards.html>
- <https://sigmodrecord.org/publications/sigmodRecord/1209/pdfs/07.industry.kulkarni.pdf>

OTHER VENDORS

- <https://illuminatedcomputing.com/posts/2019/08/sql2011-survey/>
- <https://mariadb.com/kb/en/library/system-versioned-tables/>
- https://docs.oracle.com/database/121/ADFNS/adfns_flashback.htm#ADFNS610
- https://docs.oracle.com/database/121/ADFNS/adfns_design.htm#ADFNS967
- <https://docs.microsoft.com/en-us/sql/relational-databases/tables/temporal-tables?view=sql-server-2017>
- https://www.ibm.com/support/knowledgecenter/en/SSEPGG_10.1.0/com.ibm.db2.luw.admin.dbobj.doc/doc/t0058926.html

PATCHES

- <https://commitfest.postgresql.org/48/4308/>
- <https://www.postgresql-archive.org/PROPOSAL-Temporal-query-processing-with-range-types-tt5913058.html>
- <https://www.postgresql-archive.org/SQL-2011-PERIODS-vs-Postgres-Ranges-tt6055264.html>

TOOLS

- <https://github.com/xocolatl/periods>
- https://github.com/hettie-d/pg_bitemporal
- https://github.com/arkhipov/temporal_tables
- <https://www.youtube.com/watch?v=TRgni5q0YM8>
- <https://github.com/ifad/chronomodel>

ARISTOTLE

- Photo from University of Glasgow Library, <https://www.flickr.com/photos/uofglibrary/18242587063/in/photostream/>

THANKS!

<https://github.com/pjungwir/temporal-databases-theory-and-postgres-2024>

